

Brand AI-Readiness Audit

Architecture v2 — revised after adversarial council review

Status: **proposed, ready to build** · Supersedes v1 · 2026-09-01 · 9 days to submission (2026-09-10)

1. Executive summary

Architecture v1 was submitted to a five-member adversarial council — web-standards, systems engineering, detection science, an adversarial grader, and engagement measurement — in two rounds: independent attack, then convergence on the conflicts between them. The review found **five fatal flaws**, one of which would have made the submission unshippable, and one of which was a factual error that would have caused the flagship check to fire false positives on exactly the sophisticated sites a judge is most likely to spot-check.

The resulting design is materially different from v1 and, on the grader's re-score, materially better. The headline changes:

AREA	V1	V2
Rendering	Playwright headless browser, load-bearing	Zero browser. Pure HTTP + parse. Chrome is an optional confidence upgrade that may never create a finding
Flagship check	Raw HTML vs rendered DOM diff	Extraction-difficulty tiering — is the fact in extractable <i>text</i> , or only in a script blob?
AI-bot check	Flat list of 8 tokens → "you are invisible"	Retrieval vs training taxonomy. Only retrieval-gating agents produce findings
Severity	Derived by formula <code>f(gate, blast, conf)</code>	Explicit auditable table + printed <code>severity_rationale</code>
Code layout	Shared <code>lib/</code> at marketplace root	All code in endpoint. Sub-skills are code-free knowledge skills
Blocked findings	Suppressed as moot	status: latent + <code>blocked_by</code> naming which agents
Engagement	~20 heuristics incl. CWV, above-fold, CTA verbs	~10 statically-evidenced checks , capped at medium severity
Determinism	Asserted absolutely	Qualified and made true — quantized, stable-sorted, <code>--replay</code> over fixtures

Grader re-score after applying all changes

CRITERION	V1	V2	REMAINING CEILING
1 · Detection accuracy	6	7	Static-only caps it; no outcome labels
2 · Suggested-action quality	5	7	Only if both differentiating extras ship
3 · Output design	7	8	<code>latent / coverage / measurement_basis</code> earn it
4 · Skill-format & hygiene	4	8	Biggest mover
5 · Marketplace composition	4	7	Each sub-skill needs a standalone use case
6 · Generalization	5	6.5	Weakest; counterexample hunt is the only lever

2. The five fatal findings

2.1 **FATAL** The browser cannot ship

Measured on the target machine: Playwright's Chromium is **379 MB**; the headless shell alone is **273 MB**. The submission cap is **50 MB**. Chromium ships as already-compressed binaries, so zip recovers perhaps 40% — a floor of roughly 120–160 MB. It does not fit. The only escape, `playwright install` at run time, is a ~150 MB CDN download, which violates the "self-contained, no external service required" constraint and fails outright on an offline or proxied grader machine.

Resolution: the static HTTP + parse pipeline *is* the product. It covers Gate 1 completely, Gate 3 completely, freshness completely, and most of engagement. A browser, if detected, may only stamp `render_confirmed: true` and raise confidence on a finding the static layer already fired. It may never create, delete, or re-severity one — so the graded output is byte-identical with or without it.

2.2 **FATAL** The AI-bot list conflated training opt-out with retrieval blocking

Six of the eight user-agent tokens in v1 — `GPTBot`, `ClaudeBot`, `CCBot`, `Google-Extended`, `Applebot-Extended`, `Bytespider` — control **training and corpus use only** and have no effect on whether an assistant can cite a page. Google documents that blocking `Google-Extended` does not affect Search, and AI Overviews are served from the Googlebot index. `CCBot` is Common Crawl: a corpus, not live retrieval.

As written, v1 would fire a high-severity "you are invisible to AI" finding against a publisher who made the correct, deliberate choice to block training while permitting retrieval — a textbook rubric-1 false positive on the check v1 called its differentiator.

CLASS	AGENTS	TREATMENT
Retrieval-blocking	<code>OAI-SearchBot</code> , <code>ChatGPT-User</code> , <code>Claude-SearchBot</code> , <code>Claude-User</code> , <code>PerplexityBot</code> , <code>Perplexity-User</code> , <code>Applebot</code> , <code>Googlebot</code> , <code>Bingbot</code>	Finding. This is what gates citation
Training-only	<code>GPTBot</code> , <code>ClaudeBot</code> , <code>CCBot</code> , <code>Google-Extended</code> , <code>Applebot-Extended</code> , <code>Bytespider</code>	No severity. Neutral <code>site_profile.ai_policy</code> observation with the tradeoff stated
Hybrid / ambiguous	<code>meta-externalagent</code> , undocumented agents	Observation, flagged as uncertain

2.3 **FATAL** SKILL.md is the graded artifact and was treated as packaging

The spec states that evaluation targets "the skills' instructions, checks, logic and composition" and that no individual report is graded. Architecture v1 contained roughly 1,800 words on detection logic and JSON schema and **zero words** on what any SKILL.md would say. Combined with v1's own rule that all detection lives in deterministic scripts, the honest output was five SKILL.md files reading `"run scripts/audit.py $URL , emit the JSON."` That is a wrapper, not a skill.

Resolution: the SKILL.md bodies carry the reasoning — the gate model, the archetype decision table, the severity derivation, the named false-positive exclusions, the partial- result protocol — and `scripts/` becomes the deterministic executor of exactly those written rules. Writing time is budgeted equal to build time, and moved off the last day.

2.4 **FATAL** The survey measured prevalence and was sold as validity

A survey of unlabelled sites yields the marginal distribution of site properties and no outcome variable. "78% of sites lack `Product.offers.price`" licenses exactly one inference — *this is common* — and none about detection accuracy. Worse, percentile thresholding is circular: a threshold at the 20th percentile flags 20% of sites **by construction**,

regardless of whether any has a real problem. If a defect is near-universal it becomes invisible; if a good practice is rare, healthy sites get flagged.

Resolution: the corpus is retained but renamed and re-scoped. It is a **robustness corpus** — a crash-test harness, false-positive triage set, and fixture source for `--replay`. *No threshold may be set at a percentile.* Every threshold in the codebase carries a comment naming its mechanism, not its distribution.


2.5 **FATAL** Engagement checks were a false-positive generator

Consent banners are legally mandated under ePrivacy/GDPR, so flagging their presence is a false positive on essentially every EU-facing site. Above-fold word count fires on Apple and Stripe. English action-verb lists fail on every non-English site, and the archetype detector did not detect language. A single headless run is not a Core Web Vitals measurement — Google ranks on field data at the 75th percentile, and **INP, which replaced FID in 2024, cannot be measured synthetically at all.**

Underneath sat a framing error: **bounce is the wrong target**. GA4 deprecated it in favour of engagement rate precisely because a single-page visit that answered the question is a success. For AI-sourced traffic this matters more, not less — the assistant has already summarised the answer, so the residual click is usually *verification*, and a 15-second exit is a win.

3. Core model

The three stated problems — invisibility, disengagement, staleness — are not parallel. Appendix A describes an ordered, blocking pipeline; two further concerns hang off its end.



Gates are ordered and blocking: a failure at gate n makes findings at gate $n+1$ unobservable — reported as *latent*, never suppressed.

3.1 The flagship check — extraction-difficulty tiering

Four of five council members independently named fact-level analysis the strongest idea in v1. But it was specified as a raw-versus-rendered diff, which requires the browser that cannot ship. The convergence round resolved this by observing that **the rendered DOM was never the measurement — it was only a fact source**. Removing it does not degrade the check; correctly restated, it sharpens it.

The mechanism, precisely. Assistant fetchers run a two-stage pipeline: raw bytes, then a boilerplate-stripping extractor of the Readability/Trafilatura class that converts HTML to text before the model sees anything. That stage discards `<script>`, `<style>`, nav and footer *by construction*. Hydration payloads are therefore present in the **bytes** and absent from the **text**. Present-in-bytes is not present-to-the-assistant.

So the defect is not "the HTML is an empty shell." It is: **the fact does not survive bot-grade text extraction of the raw response**. That is testable with one HTTP GET, roughly 200 ms per page, zero binaries, and no nondeterminism — both sides of the comparison come from the same response.

TIER	WHERE THE FACT LIVES	VERDICT	RATIONALE
T0	Visible prose or a heading	PASS	Directly quotable
T1	JSON-LD / microdata / RDFa	PASS	Machine-readable by design
T2	Only inside a minified hydration blob	MEDIUM	Technically in the bytes, but no assistant reliably parses escape-chunked <code>self.__next_f</code> ; Appendix F says filler-surrounded content is dropped
T3	Absent from the raw bytes entirely	HIGH	Genuinely invisible to every non-rendering fetcher

The T2 tier is the honest core of the design. A naive competitor shipping "SPA therefore invisible" will false-positive across the entire Next.js and Nuxt web. Explicitly parsing the payload islands *before* firing is a precision advantage worth naming in the README — and where a fact is found only in `__NEXT_DATA__`, that is the single strongest evidence the design can produce.

```
Evidence: "$49/mo" occurs only in __NEXT_DATA__ (byte offset 41,208)
and og:price:amount; absent from all extracted text nodes.
GET https://example.com/pricing (no JS, UA: AIReadinessAudit/1.0)
Reproduce: curl -s https://example.com/pricing | grep -c '\$49'
```

Where expected facts come from — and why it is not circular

Sourcing expectations only from structured data would make the check strongest on the sites that need it least. Two independent sources are combined:

1. **Declared literals** — values the site asserts in JSON-LD/microdata/RDFa, OG and Twitter meta, `<title>`, sitemap and feed entries, `<noscript>`, and parsed payload islands (`__NEXT_DATA__`, `self.__next_f`, `__NUXT__`, `__APOLLO_STATE__`, `window.__INITIAL_STATE__`, any `<script type="application/json">`). Test: does the literal also reach T0? *Catches: the fact exists in bytes but not text.*
2. **Role obligations** — derived from surfaces *not under test*: the URL slug we crawled, the `<title>`, and the inbound internal anchor text that led us there. A page asserting "Pricing" is a commitment to contain a currency-and-numeral token; "Contact" to a contact method; a product page to a name and price. *Catches: the fact does not exist at all.*

Where neither applies — a pure client-side shell with no nav, title or payload — empty-shell detection covers the gap, narrowly scoped to `<body>` containing a single mount element with under 30 words of text. Near-zero false-positive rate, and complementary rather than redundant.

Declared limits

Facts fetched by XHR *after* load with no payload island are invisible to this method — a real false negative, covering pure CSR SPAs, third-party price widgets and app-injected blocks. Also lost: all layout geometry, CSS-hidden versus visible distinction, and anything requiring computed styles. These are declared in `checks_skipped[]` and `coverage`, never papered over. Same-origin API endpoints referenced in the HTML are listed so a human can check them.

4. Check inventory

4.1 Gate 1 — Access & retrieval policy

ID	CHECK	NOTES
A1	Per-agent robots.txt taxonomy	Retrieval / training / hybrid. Only retrieval yields a finding. Vendor <code>protego</code> — <code>stdlib urllib.robotparser</code> is not RFC 9309-conformant on wildcards and Allow-precedence
A2	Snippet-preview controls NEW	<code>nosnippet</code> , <code>data-nosnippet</code> , <code>max-snippet:0</code> , <code>noarchive</code> in meta and <code>X-Robots-Tag</code> . Suppresses text from AI answers <i>while leaving the page indexed</i> — a page can pass every other check and remain unquotable
A3	<code>noindex</code>	Meta and header
A4	Status & redirect chains	Chain length, loops, 302-where-301-intended
A5	Canonical correctness	Self / cross-domain / to-404 / chained
A6	Sitemap integrity	Present, referenced in robots, parses, <code>lastmod</code> sane. Reported as <code>sitemap-minus-crawl-reachable diff</code> — "orphan pages" is undetectable by a bounded crawl by definition
A7	Total-block detection	A fully blocked site must emit a critical Gate-1 finding , never zero findings. See §6.2
A8	WAF / bot-management differential	Status <i>class</i> only, never body diffs; requires reproducibility across two spaced fetches; 429/503 whitelisted as inconclusive; low confidence

4.2 Gate 2 — Extractability

ID	CHECK	NOTES
B1	Extraction-difficulty tiering	T0–T3 per expected fact. The flagship
B2	Role-obligation expectations	URL slug / title / inbound anchor → obligated fact type
B3	Payload-island parsing	Six island types; fact-shaped value extraction
B4	Empty-shell detection	Narrow scope only; fallback where no role is inferable
B5	Facts locked in non-text	Content-image <code>alt</code> coverage; PDF-only facts
B6	Summarizer-hostile content APX F	Substance-to-filler ratio; key fact positioned below boilerplate

4.3 Gate 3 — Structured data & entity

ID	CHECK	NOTES
C1	Structured-data validity, three tiers	Spec-invalid (malformed JSON / <code>@context</code> / <code>@type</code>) → rich-result-required-missing → recommended-missing. Only the first two are findings; the third goes to <code>proactive_recommendations[]</code> . Absent <code>sameAs</code> is <i>not</i> invalidity
C2	Markup-vs-text contradiction	Normalise to <code>(currencyCode, Decimal)</code> . Fire only when both sides are unambiguously single-valued, same currency, differing beyond rounding. Sale-vs-list, <code>AggregateOffer</code> , variants, tax-inclusive and locale formats (<code>1.299,00 €</code>) all route to <code>unverified</code> , capped medium
C3	Entity clarity	<code>sameAs</code> targets resolve and match; name consistency across title / OG / schema / body
C4	Archetype expectation	PDP → price + availability; LocalBusiness → hours + phone; SaaS pricing showing only "Contact us"

4.4 Freshness & corroboration

ID	CHECK	NOTES
D1	Date signals	<code>datePublished</code> / <code>dateModified</code> , <code>Last-Modified</code> , sitemap <code>lastmod</code> , feed dates
D2	Staleness indicators	Copyright year, blog cadence, deprecated version strings
D3	Link rot	Broken outbound links as a decay proxy
D4	Internal contradiction	Same fact differing across pages. Threshold-free and self-evidencing
D5	Corroboration <i>infrastructure</i>	Tier 1, always runs: absent <code>sameAs</code> , no press page, inconsistent NAP, no citable stable-facts page
D6	Cross-web verification	Tier 2, key-gated. Skipped cleanly with a reason when unconfigured

4.5 Engagement — static only, severity capped at medium

Survives at full confidence: form friction (required-field count, `<label for>` / `aria-label` association, missing `autocomplete` , wrong `type` / `inputmode`) · link-purpose quality (non-descriptive anchor ratio; same anchor text resolving to different hrefs) · accessible-name presence on interactive elements · content-image `alt` coverage · `<html lang>` (gates every lexical check) · heading order and H1 · viewport meta blocking zoom (`user-scalable=no` , WCAG 1.4.4) · archetype expectation failure · promise/payoff overlap (fires only on **zero** non-stopword overlap between title+description and body) · fact locatability.

Degraded to low confidence: tap-target sizing (inline/attribute values only) · page weight (bytes, not render cost) · dead-end internal-link count (archetype-gated).

Cut outright: colour contrast · LCP/CLS/FCP/INP · everything "above the fold" · focus visibility · CTA-is-an-image · action-verb detection · PageSpeed Insights API (external call, nondeterministic, and CrUX has no record for exactly the low-traffic sites that most need auditing).

Cost of losing contrast, stated honestly. WCAG contrast was the single most unarguable number available — spec-defined, pass/fail, no taste involved. It requires computed styles and therefore a browser. Nothing replaces it, and it must not be faked from inline styles. It is declared in `checks_skipped[]` with the reason and the remedy.

4.6 The two differentiating extras

Selected by the adversarial grader as both highest-differentiation and actually buildable. They compose into one story: *can an assistant quote you, and can it link to the quote?*

- **Quotability scoring.** For each candidate 40–80 word passage, flag unresolved anaphora ("this", "our platform", "as noted above"), a missing subject entity, and dependence on a preceding heading. Assistants quote *chunks*, not pages — a passage that cannot stand alone cannot be cited. Pure static, roughly half a day.
- **Citation surface / fact locatability.** Stable heading `id` coverage, per-fact anchors, and scroll-to-text-fragment viability. Determines whether an assistant can link to a *location* rather than a homepage.

Cut to fund these: claim-fragility mapping (needs a search API we do not have), retrieval-chunk-boundary analysis (model-specific and unfalsifiable), synthetic CWV and mobile tap-targets, and the corpus from 50 sites down to 15–20.

5. Severity: a table, not a formula

Architecture v1 proposed `severity = f(gate_position, blast_radius, confidence)` and called it the central insight. The detection-science critic showed it is a lookup table with extra steps: **gate_position** and **confidence** are **both constants per check type** — gate is fixed by which check fired, confidence by whether the check is deterministic — so `f` collapses to `table[check_type] × blast_radius`. The systems critic and the adversarial grader both conceded the point.

The gate *ordering* remains real and valuable — it just belongs as the table's structure, not as a computed function. Ship an explicit, auditable table plus a printed `severity_rationale` string showing the arithmetic. A judge can read a table.

```
base          = SEVERITY_TABLE[check_id]          # defended integer, referenced to a mechanism
                                                    # rationale in references/severity-rationale.md
× blast_radius = site_wide 1.0 | template 0.8 | single_page 0.5  # banded, never a raw ratio
cap           = measurement_basis == static_heuristic -> max(medium)
               confidence      == heuristic          -> cannot reach critical
               coverage        < floor              -> no blast-radius escalation

severity_rationale: "base=high (Gate 1, blocks retrieval for 3 of 9 agents);
                    blast=site_wide x1.0; no cap applied -> high"
```

Only `blast_radius` and `coverage` are genuine runtime terms; everything else is declared per check and reviewable. Severity may consume only **bucketed** thresholds, never raw floats — this is what makes the ordering reproducible across machines.

5.1 Latent findings, not suppression

v1 proposed suppressing downstream findings on a blocked URL. All five members rejected this. It destroys information and creates a second-audit trap: the user fixes robots.txt, re-audits, and meets a wave of findings they were never warned about — which reads as an unreliable tool. The mechanism argument is stronger still: **robots is per-user-agent**, so a page disallowed to `OAI-SearchBot` may be fully open to `Claude-User`, and user-initiated fetches reach it regardless. Suppression assumes a single crawler and is unsound on its own terms.

```
{ "id": "F-014", "status": "latent",
  "blocked_by": { "finding": "F-002",
                  "agents_blocked": ["OAI-SearchBot", "PerplexityBot"],
                  "agents_unaffected": ["Claude-User", "Googlebot"] },
  "severity": "low",                // capped while latent
  "note": "Will become active once F-002 is resolved." }
```

Latent findings are excluded from `summary.total_findings`, counted in `summary.latent_findings`, and surfaced in a report section titled *"What you will face after fixing the blockers."*

6. Engineering posture

6.1 Determinism — qualified, and made true

v1 asserted determinism absolutely and was wrong: A/B tests, conditionally-rendered consent banners, `networkidle` races, CDN variance, and machine-dependent performance floats all break it. Dropping the browser removes most of this. The remainder is handled explicitly:

- ▶ Every collection stable-sorted by an explicit key before emit
- ▶ Crawl frontier seeded deterministically — sitemap order, then lexicographic BFS
- ▶ Severity consumes bucketed thresholds only; no raw floats enter the scoring path
- ▶ `--replay <fixtures/>` re-runs the decision layer over saved responses
- ▶ `run_manifest` records per-URL content hashes

The claim we make is the true one: *"the analysis layer is deterministic given fixed inputs; network observation is not."* That qualified claim scores better on hygiene than a false absolute.

6.2 Failure modes — never a silent clean bill of health

The most damaging possible output is zero findings for a site that could not be read. Every run carries `run_status` (`complete` | `partial` | `blocked` | `failed`) and per-check `coverage`, carried into the evidence string. Blast-radius escalation is disabled below a coverage floor, so a Cloudflare 403 after page three cannot produce "12 of 12 pages affected" from a denominator of three.

CONDITION	NAMED DEGRADED OUTCOME
DNS failure / site down	<code>run_status: failed</code> ; single critical finding; no check results fabricated
Redirects to login	<code>run_status: blocked</code> ; authenticated-area boundary respected; critical Gate-1 finding
WAF challenge	<code>run_status: partial</code> ; A8 fires at low confidence; coverage recorded
Robots disallows our UA	Crawl halts; <i>reported as a finding</i> , not an error

6.3 Safety posture

v1's user-agent differential test sent `User-Agent: GPTBot`. This is impersonation, and worse, it is self-defeating: claiming a user-agent binds you to that agent's robots group, so fetching paths `GPTBot` is disallowed becomes a real guardrail breach in a competition whose spec names robots.txt compliance. Resolved: a truthful identifying UA (`AIReadinessAudit/1.0 (+repo-url)`) throughout; bot policy is determined by *parsing* robots.txt, which is free, deterministic and costs zero requests; the WAF differential uses only a generic non-Chrome UA.

Read-only throughout. No authenticated areas. Honoured crawl-delay capped at 2 s with page count reduced rather than stalling — a `Crawl-delay: 10` site would otherwise consume 400 s of a 300 s budget before a single parse.

6.4 Marketplace structure

A skill folder is the unit of distribution. A `scripts/` file importing from a marketplace-root `lib/` dies the moment someone installs that skill alone — and `skills-ref validate` still passes, because it only checks frontmatter. The resolution satisfies both portability and the "SKILL.md is the artifact" finding:

All executable code lives in the entrypoint's `scripts/` as a single package. Every other skill is a code-free knowledge skill — SKILL.md plus `references/` carrying mechanism explanations, decision tables, false-positive exclusion rules and remediation snippets. Zero cross-folder imports; every folder valid standalone.

Open risk: the grader's closing warning is that four code-free skills may read as the padding rubric 5 explicitly names. Mitigation: each non-entrypoint skill must state a use case someone would genuinely invoke it alone for — *"check whether AI assistants can reach my site"* — declared in its `description` frontmatter. If a skill cannot pass that test, merge it. Three well-argued skills beat five thin ones; the spec is explicit that a single well-built skill scores fully on this criterion.

7. Validation strategy

The design has two validation instruments with sharply different jobs. Conflating them was v1's fatal error.

INSTRUMENT	WHAT IT CAN TELL YOU	WHAT IT CANNOT
Robustness corpus 15–20 diverse sites	Does the parser survive real HTML? How many findings fire per site (target: median under 15 active)? Which check fires most, and is it right on manual review? Fixture capture for <code>--replay</code>	Anything about validity. It has no outcome variable. No threshold may be derived from it
Counterexample hunt 15–20 sites, one model, one question	Does a site we score clean nonetheless fail? Does a site we score badly nonetheless succeed? One counterexample kills a check	Any rate, any threshold, any correlation. n=20 has no statistical power and Appendix E personalization is uncontrolled

On the multi-LLM council idea

The original proposal — fan a prompt across several assistants, check ground truth, log to a sheet — was rejected in v1 for lack of statistical power and absence of API keys. The council **partially overturned that rejection**, and the reasoning is worth recording: the statistical objection was correct about a design nobody should build (four models × eight prompt classes × twenty sites yields roughly one observation per cell), but wrong about the design that is actually needed.

Reduced to **one model, one question, one binary outcome per site** — *"can an assistant with web access state this site's price correctly?"* — it becomes a half-day counterexample hunt on existing consumer accounts, requiring no API keys. Three of five members argued to keep it; notably, its own proposer argued to cut it to thirty minutes. The adversarial grader's argument carried: it is the only outcome-anchored evidence in the entire submission, and therefore *the one thing that makes criterion 6 an argument rather than an assertion*.

Disposition: half a day, on Sep 8, performed by whoever is not on the critical path. Lives in `references/` labelled as an illustrative spot-check with stated limits. It never gates a build and never sets a threshold.

7.1 What ships as visible proof of generalization

Rubric 6 is tested by construction on unseen sites, so the submission must *show* its generalization reasoning rather than assert it. Two artifacts do this and cost roughly a day: `references/thresholds.md`, giving every threshold with the mechanism that justifies it (never a percentile), and `references/prevalence.md`, reporting corpus observations explicitly framed as prevalence, not derived thresholds.

8. Nine-day plan

Two structural changes from the council's draft schedule. SKILL.md prose has been moved off the final day — it is the graded artifact and cannot absorb slippage. The schema freeze is tonight, because everything downstream depends on it.

DATE	DELIVERABLE	GATE / NOTE
Sep 1 tonight	Report schema frozen; <code>marketplace.json</code> ; SKILL.md stubs with real <code>description</code> frontmatter	Slippage here slips everything
Sep 2	Fetch layer: <code>httpx</code> , <code>protego</code> per-agent robots, sitemap, deterministic frontier, <code>run_status</code> / <code>coverage</code> , fixture recorder	
Sep 3	Finding model + emitter + severity table	HARD GATE: schema-valid report end-to-end with 3 real checks. Non-negotiable
Sep 4	Gate 1 complete + Gate 3 complete	Robots taxonomy, snippet controls, canonical, JSON-LD validity, contradiction
Sep 5	Gate 2 — tiering, payload islands, role obligations + freshness	The flagship. Highest technical risk
Sep 6	Engagement static set + prioritizer + the two extras	
Sep 7	Robustness corpus; false-positive triage; fixtures → <code>--replay</code>	Delete any check firing on hand-verified-good sites
Sep 8	SKILL.md bodies as reasoning documents (full day) + counterexample hunt in parallel	Moved earlier — this is what is graded
Sep 9	README, <code>references/</code> , zip, line-by- line rubric pass	Buffer day
Sep 10	Final sanity run; submit	

8.1 Cut list, decided in advance

Scope was assessed at roughly 2.5× capacity. Cut now, not on Sep 8: Tier-2 cross-web corroboration (`checks_skipped` instead) · mobile rendering · all performance metrics · archetype detection beyond two archetypes · claim-fragility mapping · chunk-boundary analysis · corpus reduced from 50 sites to 15–20 · `--render` built only if Sep 7 finishes clean.

9. Risk register

RISK	SEV	MITIGATION
Role inference (slug/title/anchor → obligated fact) is load-bearing for the flagship check and is itself an untested heuristic	HIGH	Measure it first, not last — the first thing checked against the corpus on Sep 7. If it mislabels roles, the best check becomes the worst false-positive source
Our vendored boilerplate extractor's strip rules differ from real fetchers'	HIGH	Pin the library, ship its fixtures, declare it as a stated assumption in SKILL.md rather than an implicit one
SKILL.md prose is the graded artifact and is written late	MED	Moved to Sep 8 with a full day; stubs with real descriptions written tonight
Four code-free skills read as padding under rubric 5	MED	Each must declare a standalone use case; merge any that cannot. Consider three skills
Engagement half is being rebuilt from scratch with no appendix support and no browser	MED	Ten evidenced checks capped at medium, plus an explicit <code>limits</code> block. Ten defensible beats twenty heuristic
Nine days, three people, zero code	MED	Sep 3 hard gate; pre-decided cut list; buffer on Sep 9

10. What the council endorsed

Four ideas survived all five critics and both rounds unchanged, and should be treated as the submission's spine:

- **Fact-level analysis over volume metrics.** The evidence is a literal string present in one artifact and absent from another — unfalsifiable, non-parametric, no threshold to calibrate.

- **Contradiction detection** (markup versus text; the same fact differing across pages). Structurally high precision for the same reason.
- **checks_skipped[] with reasons.** It distinguishes "we looked and found nothing" from "we could not look" — the exact failure mode that makes audit tools untrustworthy.
- **Deterministic scripts detect; the agent narrates and never judges.** Held unanimously, and the constraint that keeps the output reproducible.